

Solving AI

Kyra, Nischal

April 21, 2023

We consider rate based neural network model with populations of inhibitory and excitatory neurons. The dynamics of the network is given by:

$$\frac{d}{dt}\mathbf{r} = -\mathbf{r}(t) + \mathbf{W} \cdot \phi(\mathbf{r}(t)). \quad (1)$$

The choice of nonlinearity ϕ is flexible and we should probably use ReLU or Sigmoid.

In particular, we divide the population of neurons to excitatory and inhibitory neurons and denote them with a subscript: $\mathbf{r}_{E/I}$. Practically, we will split the population in a 20/80 fashion with a total of 1000 neurons.

This division also induces a division of the weight matrix \mathbf{W} into block structure and we get 4 sets of matrices. The relevant one for us is the inhibitor to excitatory synapses represented by the matrix : $\mathbf{W}_{I \rightarrow E}$.

Practically, we initialize the whole matrix randomly with positive random values, and then enforce that the $I \rightarrow E$ connection is set to ~ 0 . This is also the subset of connections that will be allowed some synaptic plasticity, while all other connections basically remain frozen.

A crucial point here though: The connections are sparse and we should only connect $\sim 5\%$ of neurons presynaptically to any given neuron. We can do this by basically setting 95% of values in each row to 0. This should be fine on average. But this gives rise to a subtlety in simulation: we only update the weights that were non zero at initialization. Meaning that the weights that were set to 0 initially don't change. We can do this practically by creating a dummy motif matrix \mathbf{J} that is 1 where there is a connection and zero everywhere else and taking element wise product with \mathbf{J} of weights during synaptic plasticity.

Finally, the synaptic plasticity dynamics a la' Spike-and-Scale is as follows for N excitatory neurons and D inhibitory neurons:

$$\Delta \mathbf{W}_{I \rightarrow E} = \eta \left(\mathbf{r}_E \cdot \mathbf{r}_I^T - \mathbf{r}_E \cdot \rho \right). \quad (2)$$

where η is the learning rate and ρ is the desired firing rate written as a matrix of size $1 \times D$.

We should probably update the weights every 10th step of dynamics update, and run simulation to weight saturation. I also, think ρ should be rough equal to the \mathbf{r}_{max} at initialization. Unsure about η though, maybe should be $0.1 \times \langle r_0 \rangle$.

What do you think? Hopefully this is enough to get started.